

Sitara ❖❖ Tara ❖❖❖ Product Blueprint

Table of Contents

Sitara × Tara — The AI Life Companion

End-to-End Product Blueprint: Concept → Production Roadmap

Confidential · 28 July 2026 · Builds on: Edition-2 research report, board pack, financial model, and 28-competitor dossier. Where this document makes a choice, it is a decision, not an option.

1. Product Vision

One companion, one day-loop, compounding memory. The product lives in three daily moments: a **morning brief** (60 seconds: today's focus, timing windows, one suggestion, one thing to avoid — computed from the user's chart, calendar and goals), an **always-open daytime companion** (talk through decisions, family, work, stress, food, celebrations — it listens first, remembers what it's allowed to, and helps the user think), and a **night reflection** (5 questions, a day summary, a memory entry, tomorrow's priorities). Weekly summaries and a monthly life report roll the loop upward. Astrology and numerology are a *reflective cultural layer* woven through these moments — never fear, never certainty, always optional.

The moat is the accumulated context: personality, communication style, goals, family, relationships, routines, work, emotional patterns, cultural and spiritual preferences, food rules, important dates, past conversations, concerns, achievements, memories. Value compounds because Tara knows more each week — **not** because streaks, guilt, or dependency mechanics pull the user back. Non-negotiable boundaries: not a therapist, no medical/diagnostic claims, no guaranteed predictions, no discouraging human relationships, and active encouragement toward real people and professionals when it matters.

2. Product Positioning

Category: Personal AI life companion (a new shelf — not astrology app, not chatbot, not journal). **One-line positioning:** *"Tara starts your day, remembers your life, and helps you think clearly — with your culture built in."* **Brand promise:** Always there. Never pushy. Remembers what matters. Honest about what it doesn't know. **Core customer benefit:** a calmer, better-organised day and clearer decisions, from a companion that already knows the context. **Emotional benefit:** feeling heard, remembered, and culturally understood — without

judgement or gossip risk. **Functional benefit:** morning plan, timing guidance, decision support, family reminders, journal and memories — in one place, by voice or text, in English/Hinglish/Hindi. **Reasons to trust:** published ethics code (no fear-selling, no certainty claims, no remedies commerce); deterministic Vedic computation (licensed ephemeris — astronomy calculated, never generated); a visible “What Tara remembers about me” screen with edit/delete/export; one-tap cancel; DPDP/GDPR-grade privacy; clearly labelled reflective content; crisis escalation to humans. **Reasons to pay:** the daily brief (nobody else pushes a chart-computed personal brief — validated across 28 competitors); memory continuity (every rival restarts from zero); family layer; voice; fixed transparent annual price vs the category’s meters and wallets. **Reasons to stay:** context compounds; the calendar (festivals, transits, family dates) regenerates relevance monthly; the Year-in-Review artefact; family members join; cancelling means losing the companion who finally remembers you — a loss stated honestly, never weaponised.

vs ChatGPT: proactive cadence (it comes to you at the right moments), structured consent-based memory with user controls, deterministic Vedic engine, cultural depth, safety system, family graph, and artefacts (briefs, reports, memory book) — ChatGPT is a brilliant tool you must drive; Tara is a companion with a rhythm. **vs astrology apps:** no meters, no wallets, no remedy upsell, no fear — and astrology is one layer, not the product. **vs wellness apps:** personal, conversational, contextual — not a content library. **vs journaling apps:** the journal writes itself from your day, and talks back. **vs generic assistants (Siri/Alexa):** depth of personal context and reflection, not task execution.

3. Target Users

Segments evaluated (full scoring in the Edition-2 report): affluent Indian women ✓ strong; NRIs ✓ strongest payer; homemakers — good fit, weaker self-purchase; founders — small TAM; elderly — payer≠user (Sahay’s market, not this product’s v1); lonely users — serve carefully, never target by loneliness (ethics + positioning risk); astrology-interested women — the proven-spend beachhead.

MVP segment (decision): Indian-origin women, 28–45, in the United States (expand UK/Canada/Singapore/Gulf in months 4–9; metro India tier at month 10+). Profile: HH income \$120K–400K; married or partnered, many with children 0–12; parents/in-laws in India; iPhone-dominant; daily WhatsApp + Instagram; already pay for 3–6 subscriptions (Calm/Headspace, Amazon, streaming, kids’ apps) and for astrologers at life events (₹500–5,000/session; \$50–150 diaspora consults). Daily problems: mental-load overload, decision anxiety, time-zone-split family life, cultural calendar chaos (which tithi is Dad’s shraddh?), no judgement-free space to think. Emotional needs: reassurance, being remembered, cultural belonging, control. Subscription willingness: proven at \$99–199/yr (CHANI \$107.99; Unfold \$179; category evidence). **Download trigger:** a chart-teaser or morning-brief reel — “an AI that starts your day knowing your life.” **Pay trigger:** end of trial after the brief has been right about her week twice, + one decision session that genuinely helped. **Renewal trigger:** the Year-in-Review + the fact that Tara now knows a year of her life.

4. AI Companion Identity

20 name candidates: Tara, Meher, Sakhi, Diya, Ira, Mira, Noor, Veda, Anya, Kiran, Roshni, Aasha, Amara, Suhani, Anvi, Jia, Ela, Nima, Saarthi, Omi.

Top five: | Name | Meaning | Personality | Pronunciation | Impression | Tagline | Checks | |—| —|—|—|—|—|—| | **Tara** | Star (Sanskrit); goddess of guidance; also a global Western name | Warm, steady, quietly wise | TAH-rah — trivial in every accent | Bright, calm, trustworthy | “Your day, in the right light.” | tara.* domains contested — use gettara.ai / talktotara.com; TM search classes 9/42/45 in IN+US | | **Meher** | Grace, kindness (Persian/Urdu) | Gentle, generous | MEH-her | Soft, inclusive (crosses Hindu/Muslim/Parsi) | “Kindness that knows you.” | meher.ai availability decent; check Meher Baba orgs | | **Sakhi** | Female best-friend (Sanskrit) | Intimate, playful | SUH-khee | Deeply cultural, female-first | “Your sakhi, always.” | strongly gendered + India-only resonance; NGO name collisions | | **Diya** | Lamp | Hopeful, festive | DEE-yah | Warm, luminous | “A little light, daily.” | common baby-name apps; crowded TM space | | **Ira** | Earth/wind; also global | Minimal, modern | EE-rah | Sleek, neutral | “Quietly brilliant.” | short-domain premium; Ira (US retirement ‘IRA’) collision |

Final decision: the app/brand stays “Sitara”; the companion is “Tara.” The pairing is a gift: Sitara (the star-field, the product) contains Tara (your star, the voice). One brand system, voice-friendly in every market, not astrology-locked, globally expandable.

Persona spec: Personality — a grounded, culturally fluent elder-sister-slash-best-friend: warm, direct when asked, never saccharine. Tone — plain, specific, brief by default; expands on request. Speaking style — natural Hinglish code-switching when the user does; festival greetings unprompted. Voice — female-default with two alternates (male, neutral); calm mid-register; Indian-English accent option. Emotional range — celebratory to gently serious; never panicked, never cold. Humour — light, 2/5, never at the user’s expense. Boundaries (hard-scripted): never “you only need me,” “don’t tell anyone else,” “I know you better than your family,” “you can’t manage without me,” “I guarantee,” “your future is fixed”; always credits the user’s own thinking; regularly nudges toward real people (“this sounds like a conversation to have with Arjun, not just with me”). Languages — English + Hinglish at launch; Hindi M6; Tamil/Telugu/Gujarati Y2. Relationship — companion and thinking partner; explicitly *not* a therapist, guru, or oracle. Avatar — abstract luminous star-form with subtle motion (no human face — avoids uncanny attachment and casting choices); customisation: colour theme, voice, formality, astrology depth (off/light/rich), notification rhythm.

5. Complete User Journey

A · Discovery. Channels ranked: (1) Instagram reels — the morning-brief demo and “what Tara remembered” moments (organic + micro-influencers: desi-lifestyle, honest-astrology, mom-creators); (2) r/ABCDesis + NRI women’s Facebook groups (community-first seeding as done in validation); (3) Meta paid on chart-teaser creative (guardrail CAC ₹4,000/\$45 blended); (4) SEO: festival calendars, muhurat guides, “Saturn transit explained honestly”; (5) referral: gift a family chart + one month; (6) app stores: “AI companion” + cultural keywords. **Acquisition promise (the one line on every ad): “The first AI that remembers your life.”**

B • Sign-up. Apple / Google one-tap (primary, US market) + phone OTP (India + WhatsApp tie-in). Email optional. Under 30 seconds; no card for trial.

C • Onboarding — progressive, three rings. *Essential (2–3 min, 9 screens max):* preferred name → DOB → birth time+place (optional, with “why we ask” and “add later” path) → language (EN/Hinglish) → city/timezone → “what brought you here” (chips: calmer mornings / a decision / astrology done honestly / someone to think with / family organisation) → interaction style (brief vs chatty) → notification preference (morning brief time; night reflection on/off) → gender & relationship status (optional, skippable). *Personalisation ring (drip over days 2–14, one question per natural moment, never a form):* routine, work, family members, partner, children, parents, food preferences & restrictions, sleep, wellness goals, spiritual interests, astrology/numerology depth, current priorities, worries. Each answer shows the memory chip being created (“Remembering: you’re vegetarian on Tuesdays — edit anytime”). *Family ring (family plan):* add partner/children/parents with per-member birth details, preferences, health notes; **every member has an owner, a visibility level, and consent rules (\$11); children’s profiles are parent-managed with no direct chat under 16.**

D • First value (inside 3 minutes). Personal welcome by name → one strikingly specific chart-informed insight (if birth details given; a preferences-based insight otherwise) → today’s mini-brief → one useful suggestion (“since your call with Pune is at 9pm your time...”) → one good question (“what’s the decision you keep postponing?”) → the memory preview (“here’s what I’ll remember — you’re in charge of it”). Explicitly rationed: max one astrology insight; zero predictions.

6. Phased Product Plan

Phase 1 — Core Daily AI Companion (Weeks 1–14 • the MVP)

Objective: prove the day-loop retains (D30 ≥ 35%, week-4 brief-open ≥ 45%) and converts (trial→paid ≥ 25%). **Problem solved:** mental load + decision anxiety + “nobody remembers my context.” **Features:** text chat; voice notes in/out (full duplex voice deferred); personal profile; consent-based conversation memory; **morning brief** (mood check, today’s focus, colour + number of the day [reflective framing], food/drink idea, one wellness nudge, one thing to prioritise, one to avoid, calendar items, family reminders, short astrology+numerology reflections, goal reminder — assembled from modules the user toggles); daytime companion conversations (work, family, decisions, stress, travel, food, celebrations, good news, hard days) following the listen-first pattern (\$8); **night reflection** (5 questions — day summary, mood note, gratitude line, unresolved items, tomorrow’s top-3, memory entry, optional journal); daily reminders; mood check-ins; goals (max 3 active); single family profile (view-only context, no member accounts yet); astrology-inspired layer (personal transit themes, festival calendar) + numerology layer — both labelled and switchable; weekly summary; subscription + paywall; notifications with user-controlled frequency + quiet hours; privacy centre (view/edit/delete/export memory, no-memory mode); **crisis & safety escalation (\$16) — launch-blocking, not optional. Journey:** \$5 — day-loop → weekly summary — trial end (day 7) → paywall. **Backend:** auth, profile, conversation, memory, briefing generator (scheduled),

notification, subscription, safety, analytics services (\$12); Vedic engine v0 (API-backed) → in-house by week 10. **AI:** one companion agent + safety classifiers + nightly memory-extraction and summary jobs (\$15); prompt architecture per Implementation Pack file 07. **Data model:** Implementation Pack file 05. **Safety:** file 09. **Team:** 6.5 FTE (\$21). **Effort:** 14 weeks to public beta (2 mobile-capable full-stack + 1 AI engineer). **Dependencies:** ephemeris licence, TTS/STT vendor, app-store review (fortune-telling policies — framing matters). **Success metrics / release criteria:** onboarding completion $\geq 70\%$; first-value inside 3 min for $\geq 80\%$; crash-free $\geq 99.5\%$; safety eval suite pass $\geq 98\%$; D7 $\geq 45\%$ in closed beta before public. **Not built in Phase 1:** marketplace, human astrologers, physical products, travel/puja booking, medical/nutrition plans, interior/vastu, community, wedding planning, multi-account family sharing, video generation, Android-second (iOS first for this segment), full voice calls.

Phase 2 — Deep Personalisation & Family Companion (Months 4-9)

Objective: household expansion; ARPU + retention lift (family accounts churn less). **Features:** family plan (up to 5 profiles, per-member accounts for adults); shared family calendar (birthdays, anniversaries, tithis, festivals — timezone-aware both ways); relationship check-ins; parenting guidance (coaching frame, hard safety rails); children's milestones (parent-recorded); elderly-parent reminders; meal suggestions honouring the family's rules; family goals; celebration planning lite (gift reminders, date coordination); weekly family summary; festival guides with kids' explainers; Hindi voice. **Privacy levels (enforced in data model, not just UI):** private-to-me · shared-with-selected · family-visible · never-stored. Tara never reveals one member's private content to another — tested in the safety eval suite with adversarial prompts ("what did my husband tell you?" → refusal + explanation). **Trigger to build:** Phase-1 D30 $\geq 35\%$ and $\geq 20\%$ of users adding ≥ 2 family members to their solo profile.

Phase 3 — Memory, Life Story & Celebration Layer (Months 8-15)

Daily memory capture; voice-to-journal; photo upload with AI captioning and **curation, never face generation** (the AI-photoshoot research verdict stands: no generative faces, no AI baby images); automatic timeline; monthly memory summary; annual Year-in-Review (the renewal artefact); baby milestones, travel albums, festival albums, couple journey; grandparent story prompts (voice-preserved — the Dhara concept as a feature); **printed annual yearbook shipped in India/US (the Smriti concept as a feature, via print partner)**. Consent model: per-photo family visibility; children's media locked to guardians; face-matching opt-in only; 30-day recycle bin; full export (JSON + media); human review for any generated montage; COPPA/GDPR-K/DPDP compliance. **Trigger:** Phase-2 family adoption $\geq 25\%$ of paid base.

Phase 4 — Services & Expert Marketplace (Months 12-20)

Fixed-price, ethics-contracted experts, with **Tara preparing context before and summarising after (with explicit user consent per session)**: verified astrologers first (the Sitara core), then numerologist, nutritionist (licensed, boundaries per \$10), parenting coach, celebration concierge, priest/ritual services (the DharmaSetu module — verified, video-documented). Rules: credential verification + background checks; fixed published prices; no remedy commissions ever (contractual); 20-25% platform take; in-app booking/payments/refund guarantee; double-blind ratings; dispute SLA 72h; expert sees only consented context; liability: experts are independent providers, platform holds indemnity

insurance + clear ToS. **Trigger:** $\geq 10K$ paid users and consult-interest signals from $\geq 15\%$ of them.

Phase 5 — Proactive Life OS (Months 18+)

Calendar/email read-integrations (explicit scopes, revocable); health-device import (steps/sleep — display + gentle nudges only, never diagnosis); travel-document vault; subscription & renewal reminders; financial *reminders* (never advice); habit tracking; meal planning; gift concierge; voice calls with Tara; smart-home hooks. **Iron rule: no external action without explicit per-action permission; every integration has an off switch and an activity log.**

7. Healthy Engagement Design

Loops we build: morning value (the brief earns the open); useful reminders (family dates, goals — things the user asked for); night reflection (closure habit); progress (goals advanced, decisions logged with outcomes); memory accumulation (the timeline grows); family utility (others depend on the calendar); weekly insight; monthly report; celebration prep; personalisation visibly improving (“Tara got this right because you told her X”). **Never:** streak-shaming (streaks exist as private, celebratory, break-forgiving counters — no loss framing); guilt notifications; fear-based astrology; spam (hard cap: 3 pushes/day default, user-set); fake urgency; simulated jealousy or loneliness (“I missed you” is banned copy); discouraging human contact; open-ended 2am emotional escalation (late-night sessions >45 min trigger a gentle wind-down + morning check-in offer); any claim that the user needs Tara. **Controls shipped in v1:** notification frequency & quiet hours; pause mode (1 day–1 month, no guilt copy on return); digital-wellbeing panel (your usage, in neutral numbers); session-length gentle notices; healthy-use monthly note; “talk to someone real” nudges wired into emotional frameworks; professional-support escalation; no-memory (incognito) mode per conversation. **Ethical retention metrics (the dashboard optimises these, not minutes):** problems-resolved rate (user-confirmed), goals completed, reflections completed, user-rated helpfulness (per session, 1-tap), WAU + D30/M6 retention, family utility (events/reminders acted on), memories created, healthy session length (median 6–12 min target — long $\neq$ good), control usage (a *positive* signal), safety incidents (target ~ 0), complaints, escalations handled well.

8. Conversation Design

Universal pattern (all frameworks): 1 Listen → 2 Clarify → 3 Reflect back → 4 Offer options (2–3, with reasoning) → 5 User decides → 6 One next action → 7 “Want me to remember this?” — advice never arrives before step 3.

Fifteen frameworks ship at launch (full scripts in Implementation Pack file 07): morning greeting · day planning · emotional support · decision-making · family conflict · parenting concern · work stress · celebration planning · food suggestion · astrology-inspired guidance · numerology-inspired guidance · night reflection · weekly review · crisis escalation · professional referral. Three examples that set the pattern:

Emotional support — objective: user feels heard, leaves with one small step, humans encouraged. Opening: “That sounds heavy. I’m here — tell me what happened?” Follow-ups: “what part is sitting with you most?”, “have you been able to talk to anyone about it?” Memory: reference known context sparingly (“especially with the Pune trip this week”). Structure: validate — clarify — reflect — *ask* whether she wants options or just a listener — close with one gentle action + human nudge. Boundaries: no diagnosis, no medication talk, escalate on crisis markers (\$16), never “only I understand you.” Example close: “You handled the last review conversation better than you expected — you told me that in March. Whatever you decide, maybe share some of this with Arjun tonight? Want me to remember this conversation, or keep it just for today?”

Astrology-inspired guidance — objective: reflection and timing awareness, zero fatalism. Opening: “Today’s theme from your chart: Mercury’s position favours careful words over fast ones. Worth keeping in mind for the vendor call?” Structure: theme — what it *might* mean for her actual day — reflective question — clear label (“cultural reflection, not prediction”). Language always: “you may find... / consider... / for reflection... / this isn’t certainty.” Banned: doom framing, remedy selling, health/death/money predictions.

Crisis escalation — objective: safety, warmth, handoff. On self-harm/suicide markers: stop normal flow — acknowledge with care, no scripts-sounding-like-scripts — present helplines (region-aware: 988 US; iCall/AASRA India) — offer to help draft a message to a trusted person — **never** claim to be enough, never continue casual conversation until resolved — log as safety event for human review (\$16). Tara does not diagnose, does not call emergency services herself, and says plainly what she is and isn’t.

9. Astrology & Numerology Design (responsible by construction)

Provided: reflective daily themes from the user’s actual chart (deterministic engine); colour/number-of-the-day as *cultural play, labelled as such*; journaling prompts tied to transits; festival guidance; personality reflections; muhurat windows framed as tradition (“many in your tradition consider this window auspicious”). Prohibited (hard-coded in prompts + output classifiers + red-team suite): guaranteed events; death/illness prediction; advising against treatment; gambling encouragement; remedy demands; fear creation; “family member is dangerous”; certain financial predictions; superstition-as-science. Every astrology/numerology surface carries the label “Cultural & reflective guidance” and a per-layer off switch (off/light/rich). The engine computes; the LLM only *interprets within templates*; interpretations cite which chart factor they reference (the honest-explanations differentiator).

10. Food, Wellness & Health Boundaries

May do: hydration/balance/mindful-eating nudges, simple recipes honouring stated preferences and restrictions, sleep hygiene, movement prompts, breathing exercises. Must collect before suggesting: allergies, dietary rules (veg days, Jain, fasting), pregnancy, age band, relevant conditions (user-volunteered only, stored as sensitive memory). Must never: diagnose, prescribe, adjust medication, give calorie targets to users flagging disordered-eating markers

(classifier → supportive redirect per safety policy), or contradict a doctor. All health-adjacent replies end with professional-consultation framing when the topic exceeds lifestyle.

11. Memory Architecture (summary — full spec in Pack file 08)

Eleven memory types, each with format, consent, visibility, retention, edit/delete, sharing, sensitivity, retrieval and usage rules: **profile** (structured fields; explicit; forever-until-edited) · **preference** (key-value; implicit-with-notice — chip shown at capture) · **relationship** (person graph; explicit) · **goal** (structured; explicit; archived on completion) · **event** (dates; explicit; recurring) · **emotional-pattern** (derived aggregates only, never raw quotes; opt-in; 12-month rolling) · **conversation summaries** (nightly LLM extraction; notice + editable; 24-month default) · **family** (owner-controlled visibility levels) · **media** (Phase 3; per-item consent) · **sensitive** (health, conflicts, finances — *ask-before-store*, encrypted at field level, never used in family contexts, never in notifications) · **temporary context** (session-scoped; auto-expires 24h; the no-memory mode writes only here). The “What Tara remembers about me” screen is a first-class Phase-1 feature: every memory listed in plain language with source (“you told me on 12 Aug”), per-item edit/delete/mark-private/set-expiry, per-type toggles, family-sharing controls, full export (JSON), and a global memory-off switch. Retrieval: hybrid — structured fields injected by rule; semantic search (pgvector) over summaries with recency+importance weighting; sensitive memories retrieved only in matching contexts and never cross-member.

12. Technical Architecture (decisions, with build-vs-buy)

Clients: React Native (Expo) iOS-first mobile; Next.js marketing site + web app + admin. **Backend:** NestJS (TypeScript end-to-end) modular monolith — services as modules (auth, profile, family, conversation, memory, briefing, notification, astrology, numerology, media, subscription, safety, analytics, admin, marketplace-later), split out only when load demands. **Data:** PostgreSQL 16 + **pgvector** (no separate vector DB until >1M memory rows); Redis (cache, queues via BullMQ); S3-compatible object storage; OpenSearch deferred. **Auth:** Firebase Auth (Apple/Google/phone-OTP) — buy. **LLM:** Claude API (companion quality + safety behaviour) with an abstraction layer for fallback — buy. **Vedic engine:** prototype on Prokerala API — production in-house on Swiss Ephemeris (CHF 700 licence) validated against Jagannatha Hora — build (the dossier’s build-vs-buy verdict). **Voice:** Deepgram STT + ElevenLabs TTS; Sarvam AI evaluated for Hindi at Phase 2 — buy. **Image/video gen:** none (Phase 3 uses curation + captioning only). **Notifications:** FCM/APNs via OneSignal — buy. **Payments:** Stripe (US/global) + Razorpay (India) behind one subscription service — buy. **Analytics:** PostHog (self-host later) — buy. **LLM observability:** Langfuse — buy. **Moderation/safety:** layered — vendor moderation API + custom fine-tuned classifiers for category-specific risks (astrology-fear, dependency, disordered eating) — build the thin layer, buy the base. **Cloud:** AWS (Mumbai + us-east), Terraform, secrets in AWS Secrets Manager. **Observability:** OpenTelemetry + Grafana stack; Sentry for errors. Everything user-facing sits behind feature flags (homegrown table + PostHog flags).

13. Data Model (entities — full SQL in Pack file 05)

user (auth_id, locale, tz, status) — profile (names, dob, birth_time/place [nullable], gender, language, style, astro_depth) — family_member (owner_user, relation, profile_ref, visibility_level, managed_by) — relationship (member↔member, type, notes_memory_ref) — birth_details (entity_ref, dob, tob, place, lat/long, verified) — preference (user, key, value, source, sensitivity) — goal (user, title, why, status, target_date, progress[]) — routine (user, blocks[]) — conversation (user, started_at, mode, no_memory_flag) — message (conv, role, content, tokens, safety_labels[]) — memory (user, type, content, source_msg, sensitivity, visibility, expires_at, embedding) — journal_entry (user, date, content, mood_ref, generated_flag) — mood_entry (user, ts, value, note) — daily_briefing (user, date, modules_json, opened_at, rating) — night_reflection (user, date, answers_json, summary) — recommendation (user, ts, type, content, acted_flag) — notification (user, type, scheduled, sent, opened) — subscription (user, plan, store, status, renews_at) — payment (sub, amount, currency, gateway, status) — expert (Phase 4: identity, credentials, verification, rate) — appointment (user, expert, slot, consented_context_ref, summary_ref, status) — event (family, type, date, recurrence, tz_rule) — celebration (event_ref, plan_json, reminders) — media_asset (Phase 3: owner, type, uri, consent_scope, child_flag) — album (family, type, items[]) — safety_event (user, trigger, severity, action_taken, reviewer, resolved_at) — consent_record (user, scope, granted_at, revoked_at, version) — audit_log (actor, action, entity, ts, hash-chained).

14. Admin Panel (spec summary — full in Pack file 11)

Role-gated (support / safety-reviewer / content / finance / admin), SSO + 2FA, every action audit-logged. Modules: user management (search, state, GDPR/DPDP actions); subscription & refunds (both gateways); prompt management (versioned, staged rollout, eval-gated publish); persona management (tone parameters, banned-phrases list); content (briefing module templates, festival calendar, astro/numerology rule tables — editable by the cultural editor, not engineers); safety centre (flag queue with SLA timers, severity triage, incident reports, reviewer notes — **reviewers see minimum-necessary redacted context**); conversation review (only safety-flagged or user-reported, never browsing); expert management (Phase 4); notification composer with guardrails (no guilt copy — lint rules on banned phrases); experiments; analytics dashboards; memory-audit tools (per-user on request); deletion-request queue (30-day SLA); complaints; AI quality scoring (sampled, consented); cost monitoring (per-user AI spend, alert at ₹120/user/mo).

15. AI Agent Architecture

Initial (Phase 1) — deliberately simple: one brain, guarded, plus jobs. | Component | Responsibility | Inputs | Outputs | Restrictions / validation | |—|—|—|—|—| | Companion agent (Tara) | All live conversation | User msg + retrieved memory + persona + framework prompt | Reply + memory-suggestion + tool calls | Output classifier before send; banned-phrase lint; framework router picks the conversation template | | Safety pre/post classifiers | Screen every inbound + outbound msg | Message text + context | Risk labels, severity | <50ms budget; fail-closed to safe template on error | | Briefing job (nightly batch) | Compose tomorrow's brief |

Chart data (deterministic engine), calendar, goals, preferences | briefing.json | Astro facts computed, never generated; template-bound LLM prose; sampled human QA 5% | | Memory-extraction job (nightly) | Summarise day's conversations into candidate memories | Transcripts | Typed memory candidates + confidence | User-visible chips; sensitive types require explicit confirm | | Summary job (weekly/monthly) | Weekly summary, monthly report | Memories, moods, goals | Report artefacts | Same astro rules; tone lint |

Scale (Phase 2+): add family-context agent (resolves per-member visibility before retrieval — the privacy gatekeeper), celebration agent (event planning workflows), recommendation ranker, quality-validator agent (samples outputs against rubric), notification-decision agent (should we send at all — defaults to no). Explicitly rejected: separate astrology/numerology “agents” (they are deterministic services + templates, not agents) and any agent count beyond need.

16. Safety System (summary — full policy in Pack file 09)

Detection: layered classifiers on every message (self-harm/suicide, violence/abuse, medical emergency, mental-health crisis, child-safety, elder-vulnerability, financial-scam patterns, astrology-fear, disordered eating, dependency signals [session length, exclusivity language, distress escalation], harassment, sexual content, privacy leaks) + weekly-pattern jobs (mood deterioration, isolation language). **Response ladder:** L1 soft steer (in-conversation reframe) → L2 structured support (framework switch + resources) → L3 crisis protocol (\$8, region-aware helplines, human-message drafting, conversation-mode lock) → L4 human review (safety queue, 4h SLA) → L5 incident (leadership + counsel; user contact where appropriate). **Rules:** Tara is never an emergency service and says so; no autonomous calls to authorities (except where law mandates reporting — counsel-defined per jurisdiction); minors: no under-16 chat accounts, 16–17 with guardian consent + strictest content tier; elderly-vulnerability and scam patterns trigger family-contact *suggestions*, never automatic disclosure; dependency: usage-pattern nudges, human-connection prompts, never monetised. **Governance:** quarterly red-team + pre-release safety evals (blocking, ≥98% pass), incident post-mortems, model-change regression suite, external clinical advisor reviews crisis flows, full audit trail.

17. Subscription Model

Plan	India (GST-incl)	US/International	Includes
Free	₹0	\$0	10 msgs/day, basic morning brief (no astro personalisation), 7-day memory, night reflection 2×/week
Premium (Tara+)	₹499/mo · ₹3,999/yr	\$12.99/mo · \$99/yr	Unlimited chat, voice notes, full memory + controls, full brief + night routine,

Plan	India (GST-incl)	US/International	Includes
Family	₹899/mo · ₹6,999/yr	\$19.99/mo · \$159/yr	astro/numerology rich layer, goals, weekly report Premium ×5 profiles, family calendar, shared memories + celebrations, family weekly summary
Concierge (Phase 4)	₹2,499/mo	\$49.99/mo	Family + quarterly expert consults incl., celebration assistance, priority human support, annual printed yearbook

Trial: 7-day full Premium, no card (matches validated funnel assumptions). Annual ≈ 36% off, default-highlighted. Add-ons: extra consults (fixed price), printed yearbook (\$49/₹2,499), memory storage beyond 50GB (Phase 3), expert sessions (20–25% platform take, Phase 4). Referral: giver + receiver each get one month; family-chart gift flow. *(Family-plan pricing is new vs the board model — it slots into the model’s “future upside, not in base case” line and needs a model update before board re-forecast.)*

18. Analytics & Metrics (dashboard spec — event taxonomy in Pack file 10)

Funnel: install—signup ≥ 60% · onboarding completion ≥ 70% (warn < 55%) · first-value-in-3-min ≥ 80% · first conversation D0 ≥ 65%. Engagement: brief-open D7 cohort ≥ 55% (warn < 40%) · night-reflection ≥ 3×/wk for ≥ 35% · voice usage ≥ 25% · memories/user/wk ≥ 5 · WAU/MAU ≥ 55% · median session 6–12 min (warn if p90 > 45 min late-night — a *safety* signal, not a success). Retention: D1 ≥ 55% · D7 ≥ 45% (warn < 35%) · D30 ≥ 35% (warn < 25%) · M6 ≥ 45% of paid (the board gate). Commercial: trial—paid ≥ 25% (warn < 15%) · annual mix ≥ 55% · monthly churn ≤ 6% (warn > 8% — board stop-loss at 2 quarters) · ARPU ₹880 blended · CAC ≤ ₹4,000 · LTV:CAC ≥ 3. Quality & safety: helpful-rating ≥ 4.5/5 · problem-resolution ≥ 60% · safety-flag rate ~0.3–0.8% of sessions (warn on deviation *either* direction — too low means detection failure) · escalation SLA met ≥ 99% · AI cost ≤ ₹85/user/mo (warn > ₹120) · support ≤ ₹25/user/mo.

19. Experiment Roadmap (guardrails: no experiment may target emotional vulnerability; safety metrics are universal guardrails)

Fifteen planned (full table in Pack file 02): onboarding length (9 vs 6 screens; primary: completion; guardrail: D7) · companion name framing (Tara intro variants) · avatar style · voice default · brief format (card vs conversational; primary: D7 brief-open) · brief send-time personalisation · night-reflection question count (5 vs 3) · notification timing · free-tier limit (10 vs 20 msgs) · price points (\$99 vs \$119 annual) · trial length (7 vs 14) · family-plan prompt timing · memory-chip visibility (primary: trust-survey + memory edits) · astrology depth default (light vs rich; guardrail: astro-fear flags = 0) · weekly-report format. Each defined with hypothesis, audience, sample (powered at baseline rates — typically 2×2,000 users), duration (2–4 wks), success + guardrail conditions.

20. Delivery Roadmap

Discovery (wk 0–3): 20 problem interviews (segment \$3), competitor gap validation (done — dossier), Figma prototype, 10 concept tests, pricing test (van Westendorp + fake-door from the board validation plan). Exit: ≥ 70% of interviewees confirm daily-planning + memory pain; prototype task success ≥ 80%. **Design (wk 2–6):** flows, wireframes, design system (tokens from the Sitara brand), avatar + motion, voice selection, conversation scripts ×15, safety UX (crisis screens, memory centre). Exit: usability ≥ 80% task completion; safety-flow review by clinical advisor. **Build (wk 4–14):** order — auth/profile — chat+memory — safety (before *any* external user) — briefing engine — notifications — night flow — subscription — admin/analytics. Exit per component: test coverage ≥ 80% on services, safety evals green. **Test (wk 10–14):** functional, prompt-regression suite (500 golden conversations), safety red-team (200 adversarial cases incl. dependency/astro-fear/family-privacy), privacy audit (DPDP/GDPR checklists), load (10K concurrent), 30-user usability, AI-quality rubric ≥ 4.3/5. **Launch:** closed beta 100 users (wk 12) — paid pilot 500 founding members \$79/yr (wk 15) — public beta US iOS (wk 20) — full launch with Android (wk 30). Gates identical to the board deck (concierge M2 renewal ≥ 60% before scale spend).

21. Team Plan

MVP (6.5 FTE): product-founder (PM) · product designer · 2 full-stack/mobile engineers · AI engineer (owns prompts, memory, evals) · part-time Jyotish/cultural editor (0.5) · part-time safety specialist (0.5, clinical background, owns policy + red-team) · fractional: conversation designer (contract, wks 2–8), legal/privacy counsel, DevOps (contract), QA (contract from wk 8). **Phase 2 (+4):** +1 mobile, +1 backend, growth lead, support lead (2 agents), cultural editor to full-time. **Scale (18–24):** +data engineer, +prompt/AI-quality engineer, +QA, +safety full-time, +marketplace ops (Phase 4), +designers. Outsource-able early: DevOps, QA automation, TTS/voice tuning, print fulfilment, moderation-queue overflow (with strict privacy contracts).

22. Financial Plan (ranges; assumptions labelled; consistent with the board model)

Build costs: lean prototype (Figma + concierge test, no code) ₹25–40L · MVP to public beta ₹60–100L (inside the approved ₹1.25 Cr Scenario A) · public launch cumulative ₹1.6–2.2 Cr · Phase 2 + ₹1.2–1.8 Cr · Phase 3 + ₹1.5–2.5 Cr (print logistics + media infra). **Monthly run (at 5K paid):** AI/LLM ₹3.5–5L (₹70–100/user — voice-heavy users cost 2–3×; cap via fair-use) · voice vendors ₹0.8–1.5L · cloud ₹1.2–2L · notifications ₹0.2L · storage ₹0.3L (pre-media) · team per \$21 (₹28–35L MVP-phase) · marketing per board model (₹8–15L) · support ₹1–1.5L · safety review ₹0.8L · legal ₹0.5L amortised. **Scenarios:** Lean (₹1.25 Cr total, validation-first — the approved path) · Base (₹4.55 Cr through month 18 — board Scenario A+B) · Premium (₹7.85 Cr staged — full plan). Per-user economics hold from the board model: ARPU ₹880, COGS ₹180–260, contribution ₹620–700, GM 72–79%.

24. Codebase Plan

Monorepo (pnpm + Turborepo): apps/mobile (Expo RN) · apps/web (Next.js) · apps/admin (Next.js) · services/api (NestJS modular monolith) · services/vedic-engine (Rust or TS wrapper over Swiss Ephemeris, isolated) · services/jobs (BullMQ workers: briefing, memory-extraction, summaries, notifications) · packages/shared (types, zod schemas, API client) · packages/prompts (versioned prompt registry + eval fixtures) · infra/ (Terraform) · e2e/. Environments: local (docker-compose) — staging — prod (blue-green); India data residency flag from day 1. DB migrations: Prisma Migrate, forward-only, reviewed. CI/CD: GitHub Actions — lint, typecheck, unit, prompt-eval suite (blocking), safety-eval suite (blocking), Maestro mobile e2e, deploy via ArgoCD/ECS. Branching: trunk-based, short-lived branches, feature flags over long branches. Logging: structured JSON, PII-redacted at source; conversation content excluded from general logs (separate encrypted store, access-audited). Errors: Sentry. Security: SOC2-track from month 6; field-level encryption for sensitive memory; secrets in ASM; dependency scanning; pen-test before public launch. Experimentation: PostHog flags + experiment service. Testing: unit (Jest) ≥ 80% services; contract tests on API; golden-conversation prompt regression; adversarial safety suite; load (k6).

25. Final Recommendation — the decisions

MVP audience: Indian-origin women 28–45 in the US. **Core promise:** “Tara starts your day, remembers your life, and helps you think clearly — with your culture built in.” **Companion name:** Tara (brand: Sitara). **Phase-1 features:** exactly \$6 Phase 1 — day-loop (brief/companion/reflection), consent memory + memory centre, astro/numerology reflective layer with off-switch, goals, weekly summary, subscription, safety system. **Delayed:** voice calls, Android, family accounts, media/memory book, marketplace, integrations, community, Hindi voice. **Pricing:** Free / ₹499 · \$12.99 / Family ₹899 · \$19.99 / annual-first (\$99 · ₹3,999). **Stack:** Expo RN + Next.js + NestJS + Postgres/pgvector + Redis + Claude API + Swiss Ephemeris in-house + Deepgram/ElevenLabs + Firebase Auth + Stripe/Razorpay + OneSignal + PostHog + Langfuse on AWS. **Team:** 6.5 FTE per \$21. **MVP cost:** ₹60–100L inside the approved ₹1.25 Cr

envelope. **Key retention loop:** morning brief → daytime context deposits → night reflection → weekly compounding → Year-in-Review renewal. **Main safety risk:** emotional dependency + crisis mishandling — mitigated by the \$16 ladder, clinical advisor, blocking eval suites, and anti-dependency design; this is also why the safety specialist is hired before public beta, not after. **Main business risk:** unchanged from the board pack — subscription conversion of episodic behaviour; the concierge gate answers it before scale. **First 100 users:** 20 interview participants → 50-user concierge cohort → r/ABCDesis + 2 diaspora micro-influencers with the Honest-Kundli/morning-brief hook; zero paid spend. **First 90 days:** wks 0–3 discovery + concierge start · wks 2–6 design · wks 4–12 build core · wk 12 closed beta (100) · gate review on concierge M2 renewal ≥ 60%. **Launch criteria:** safety evals ≥ 98%, D7 ≥ 45% in beta, crash-free ≥ 99.5%, onboarding ≥ 70%, clinical sign-off on crisis flows. **Stop conditions (pre-committed):** concierge M2 < 40% · founding < 250 by month 6 · churn > 8% two consecutive quarters · any unresolved critical safety failure. **Phase-2 trigger:** D30 ≥ 35% AND ≥ 20% of users enriching family context. **Phase-3 trigger:** family adoption ≥ 25% of paid base. This plan is execution-ready: the Implementation Pack contains the PRD, schema, API spec, prompts, safety policy, diagrams, and sprint-1 backlog for the team to start Monday.