

Board Pack — Slide-by-Slide Outlines (Deck 1 & Deck 2)

Project Sitara · Confidential · 28 July 2026 Conventions: ₹ lakh / ₹ crore, Indian numbering. FX assumption ₹90/USD. Facts = cited research (Edition-2 report). Benchmarks = named competitors. Assumptions/Management estimates = labelled. Projections = from the bottom-up Excel model (base case unless stated).

DECK 1 — BUSINESS IDEA & BOARD CONVICTION (30 slides)

S1 · Cover — Purpose: set premium, decisive tone. Headline: “The Reassurance Economy — Building India’s First Trust-First AI Life-Guidance Subscription.” Sub: “From 14 researched categories to one decisive recommendation: Project Sitara.” Content: Confidential — Board Presentation · 28 July 2026. Visual: dark navy cover, star-chart motif. Notes: two research phases, 21 workstreams, ~670 sources; today we ask for a validation decision, not a scale decision. Takeaway: this is a strategic conclusion, not an ideas tour.

S2 · The Decision We Need to Make — Purpose: frame the ask upfront. Headline: “Approve ₹1.25 crore to validate and build v1 of our highest-scoring opportunity.” Content: considering = entry into high-value AI subscriptions for affluent Indians/NRIs; why now = category evidence + AI capability crossover in 2025-26; asked today = Scenario-A validation approval (₹1.25 Cr) + stop-loss; NOT asked = full launch capital (Gate 2, board review month 6). Visual: 3-box decision frame. Notes: staged-capital philosophy; maximum downside is capped at the stop-loss. Takeaway: a capped, gated bet.

S3 · Executive Summary — Purpose: whole recommendation on one slide. Headline: “Proven spending, an unclaimed premium niche, one winner.” Content: opportunity = NRIs + affluent Indians already pay recurrently for reassurance/care/rituals (AstroTalk ₹1,182 Cr FY25, 80% repeat; Samarth ₹11-17K/mo; Sri Mandir NRI ARPU 10x); gap = no trust-first premium product; top 3 = Sitara 8.1 · Sahay 8.0 · KalaOS 7.6; winner = Sitara; initial investment ₹1.25 Cr; monthly break-even ~M30 (base); 3-yr revenue path ₹1.42 Cr — ₹6.34 Cr — ₹12.16 Cr (model outputs); next step = 8-week concierge validation. Visual: stat cards. Takeaway: decision-ready.

S4 · Why Now — Purpose: macro timing. Headline: “Four shifts converged in 24 months.” Content: (1) AI crossed the credibility line — AstroSage’s AI astrologers rated 4.6 vs 4.3 human at ~90% margins (fact); (2) subscription comfort — India ~119-148M paid OTT subs; NRIs pay \$10-25/mo consumer bands (facts); (3) wealth — 35.4M diaspora, \$129.4B remittances 2024, ~100M affluent Indians by 2027 (facts: MEA, World Bank, Goldman Sachs); (4) category heat — AstroTalk IPO-bound (~₹12,500 Cr talk). Visual: 4-quadrant with stats. Takeaway: the window is open and visible.

S5 · The Target Customer — Purpose: who pays. Headline: “Hard-currency income, Indian-heart spending.” Content table: NRI women 27-45 (US/UK/CA/Gulf, \$120-400K HH, WTP \$13-25/mo, need = timing/reassurance); NRI children of ageing parents 32-55 (WTP \$49-299/mo, need = parent safety); affluent metro Indian women 25-45 (₹15L+ HH, WTP ₹300-2,000/mo); devout diaspora families (WTP ~\$81/yr proven). Visual: segment cards with WTP bars. Notes:

Indian-Americans = highest-earning US ethnic group (~\$150K median). Takeaway: the payer is the diaspora wallet.

S6 • What They Already Pay For — Purpose: prove spend exists. Headline: “We are reorganising existing spend, not inventing it.” Content: astrologers ₹500-1,500/session (repeat); therapy ₹800-3,500/session; tutors \$15-40/hr; eldercare ₹11-17K/mo; pujas \$10-100+; weddings ₹39.5L avg; parenting coach \$40-200/session; premium apps \$70-200/yr. Visual: horizontal spend-band chart. Takeaway: every rupee in our model is already being spent somewhere worse.

S7 • The Real Customer Problem — Purpose: emotional core, customer language. Headline: “Every astrologer tells me something different — and none of them remembers me.” Content: verbatims from research: “paid ₹500... generic advice”; “₹15,000 puja to remove bad luck”; “This is the tax of living abroad” (eldercare); “no proof the puja happened”; “my daughter answers in English.” Needs: reassurance, continuity, trust, time, cultural belonging. Visual: quote wall. Takeaway: recurring anxiety, transactional solutions.

S8 • Why Existing Solutions Fail — Headline: “Episodic, memory-less, and often predatory.” Content: per-minute marketplaces monetise anxiety (fear-upsell documented; Trustpilot 3.3/5 category leader); no continuity between consults; Western apps = dark-pattern auto-renew (Nebula \$9.99/wk complaints); free tools = no trust, no fulfilment, no memory. Visual: failure matrix (fragmented / one-time / no-memory / low-trust × competitor types). Takeaway: the gap is trust + continuity, not information.

S9 • Why AI Changes the Economics — Headline: “AI turns a per-minute human business into an 75%+ gross-margin subscription.” Content: personalisation at scale, 24×7 across time zones, voice, memory/continuity, Hinglish/Hindi, proactive alerts, 90% margin proof (AstroSage AI); what AI cannot do alone: high-stakes trust moments (human astrologer escalation), ethics enforcement, deterministic astronomy (computed, not generated). Visual: cost-per-interaction bar: human ₹30-130/min vs AI <₹1. Takeaway: the economics flipped; humans move to the premium layer.

S10 • The Opportunity Universe — Headline: “14 categories, 20 scored ideas, one map.” Content: theme map — Family & children (Bolo, Palna); Memory & celebrations (Smriti, Dhara, photoshoots, events, birthdays, kitty); Wellness & companionship (Saathi, Mann-ki-Baat, Sahay); Spirituality & culture (Sitara, DharmaSetu); Travel & lifestyle (YatraOS, IndiaDesk); Personal/business content (KalaOS, VyaparLens, Brand Studio); Home (GharShanti). Visual: bubble map sized by score, coloured by verdict. Takeaway: breadth was examined; rigour did the choosing.

S11 • How We Evaluated — Headline: “One weighted model, applied to everything.” Content: 12 criteria with weights (WTP ×1.5, retention ×1.5, AI advantage ×1.25, market ×1.25, subscription fit ×1.25, urgency/acquisition/scalability/defensibility ×1.0, ops/trust/founder-fit ×0.75); kill-rules: replaceable-by-free-ChatGPT, one-time purchase shape, payer≠user without a channel, regulatory minefields. Visual: weight donut + kill-rule chips. Takeaway: the same yardstick for every idea — no pet ideas.

S12 • Full Opportunity Ranking — Headline: “Twenty ideas, honestly scored.” Content: ranked table (rank, idea, customer, price band, subscription strength, margin, ops, score) —

bottom 10 greyed with one-line rejection reasons (Lensa –98% novelty collapse; Evite 25-yr ceiling; gaming-law kill; one-time Vastu; free-Gemini deflation). Visual: scored bar list low–high. Takeaway: the rejections are evidence, not opinions.

S13 • The Rejected Tier (why not) — Headline: “What we will NOT build, and the evidence that spared us.” Content: AI photoshoots (novelty physics + free Google + child-safety warnings); celebrations/birthday/kitty (frequency math, regulation); Vastu (no consumer subscription exists anywhere); travel (affiliate-shaped); standalone companion (payer≠user + safety); personal-brand studio (quarterly cadence). Visual: tombstone cards with the killer stat each. Takeaway: capital preserved by research.

S14 • Contender #10-7 (rapid) — IndiaDesk 6.3 (asset-anchored retention, ops sprawl) · VyaparLens 6.4 (Dresma-proven, services-margin) · Saathi 6.4 (BetterHelp-priced, human-margin cap) · Palna 6.6 (Joy-validated model, thinnest moat vs free ChatGPT). One row each: price, why-pay, why-not-higher. Visual: 4 compact cards. Takeaway: real businesses, not board-scale winners.

S15 • Contender #6-5 — Bolo 6.9 (voice-first Hindi for diaspora kids; child-churn graveyard) · Smriti 7.1 (NEW: FamilyAlbum-30M proof + grandparent print corridor; aging-out churn caps LTV). Visual: two cards with score bars + key stat. Takeaway: strong niches; retention physics rank them below the top 4.

S16 • Contender #4 — DharmaSetu 7.4 — Headline: “The clearest WTP signal in the corpus — and a fulfilment business.” Content: 10x NRI ARPU (Sri Mandir, fact); ritual calendar = inherently recurring; verification gap = documented #1 complaint; held back by: subscription unproven (₹1/day pilot never scaled), fraud-scarred category, Sri Mandir can build it. Visual: ARPU 10x bar + risk chips. Takeaway: a top-three-adjacent idea, best as the winner’s phase-2 cross-sell.

S17 • Top Three Side-by-Side — Headline: “Three fundable ideas; one best first bet.” Content: comparison table — customer, need, price, market, frequency, retention, AI advantage, GM%, acquisition, ops, risk, 3-yr revenue potential (ranges from idea-comparison model). Visual: 3-column scorecard. Takeaway: all three clear the bar; sequencing is the decision.

S18 • Top Idea #3 — KalaOS 7.6 — What: India-priced AI OS for wedding studios (culling/editing/albums/delivery). Attractive: only idea where every WTP layer is proven by profitable companies (Imagen \$10M ARR profitable; Aftershoot 188K photographers; HoneyBook \$140M ARR); ops-light B2B; no safety minefield. Why not first: smaller ceiling (India software TAM \$50-150M est.), head-on race vs Indian-founded incumbent, B2B tools brand ≠ consumer platform ambition. When to revisit: if consumer validation fails, this is the safest pivot; or as year-3 B2B2C adjacency. Visual: proof-stack bars. Takeaway: best-proven economics, smallest dream.

S19 • Top Idea #2 — Sahay 8.0 — What: NRI parent-care command centre (voice-AI daily calls in Indian languages + family dashboard + verified on-ground partners). Attractive: rawest pain in corpus; structural retention; phase-2 de-risking (ElliQ 30-40 interactions/day sustained; families pay \$15-40/mo for AI calls — inTouch/Meela; IamFine a decade at \$14.99). Blockers: ops 6.5/10 at trust-delivering tiers; funded incumbents (Emoha ₹85 Cr); emergency liability; fails “scale without humans” test at the moment of truth. Future: the natural second

product — shares the NRI family graph; Combination-6 endorsed (calls + life-story + albums). Visual: pain/proof panel. Takeaway: the bigger company, the harder company — sequence it second.

S20 • The Winner — Sitara 8.1 — Headline: “A trust-first AI Vedic life-planning companion.” One-line: knows your chart, your family and your life story; daily timing brief; honest explanations; verified human astrologers; no fear-selling — \$99-199/yr NRI, ₹3,999-7,999/yr India. Why it wins the weighted model: proven ₹1,182 Cr category with 80% repeat; AI outperforms incumbent humans (4.6 vs 4.3) at 90% margins; fastest-growing segment unserved (international +4.2x); ops 3/10; documented trust vacuum = brand opening. Visual: winner card + 5 proof stats. Takeaway: the best combination, not the biggest single number.

S21 • Persona — “Priya” — 34, PM, New Jersey; H-1B-GC; married, one child; parents in Pune; HH \$210K; pays Calm/Peloton/Kumon; consulted astrologers at wedding + house purchase; follows 2 Instagram astrologers; finds AstroTalk “sketchy but I go back.” Trigger: job offer / kundli matching / house muhurat. Objection: “is this real Jyotish or an app gimmick?” WTP: \$199/yr in one tap if trusted. Secondary persona: Meera, 39, Bengaluru marketing head (₹28L HH) — ₹999/mo tier. Visual: persona cards. Takeaway: she already spends more than our ask.

S22 • Product Experience (journey) — Discovery (IG astrologer collab) — free Honest-Kundli teaser (magic moment: chart-specific insight in 60s) — onboarding 10-min (birth data, goals, tone) — daily 60-second brief (habit spine) — week-1: first decision session — monthly transit report — quarterly human consult (premium) — annual Year-Ahead artefact = renewal moment — refer via free family chart. Visual: journey line with moments-that-matter. Takeaway: habit by day 7, trust by month 3, family lock-in by month 6.

S23 • Product Storyboard (screens) — Home (today’s brief + timing windows); Companion chat (memory chips visible: “remembers: job search, Ananya’s school”); Decision session (structured Q&A — recommendation + why); Family (charts, festival calendar); Consult booking (verified astrologer, fixed price, context handoff); Year-Ahead report (shareable premium artefact). Visual: 6 wireframe mocks. Takeaway: an OS for life-timing, not a horoscope feed.

S24 • Why Users Pay / Why They Stay — Pay: functional (timing windows, one trusted answer), emotional (reassurance without judgment), convenience (2am across time zones), identity (modern-premium way to honour tradition), financial (< one consult/month), family (shared calendar). Stay: memory compounds — guidance sharpens; calendar auto-recurs (festivals/transits/dashas); human consult anchors trust; family charts embed; Year-Ahead resets the loop; leaving = losing the astrologer who finally remembers you. Ethics: no manipulative retention; easy cancel. Visual: two-column value/retention engine diagram. Takeaway: retention is engineered, not hoped.

S25 • Why Free ChatGPT Can’t Replace It — Layers: deterministic Vedic computation engine (never hallucinated); years-deep life log + outcomes; proactive transit/muhurat alerts; verified human panel under ethics contracts; family graph; secure DPDP-grade storage of birth data; premium artefacts; brand covenant (no fear-selling) in a low-trust category. Visual: moat stack. Notes: ₹399 ChatGPT Go is exactly why we sell continuity + humans + ethics, not text. Takeaway: the moat is memory, trust, and people — not the model.

S26 • Competitive Landscape — 2×2: Trust/ethics (y) × Continuity/personalisation (x). AstroTalk (low-y, low-x, huge), AstroSage AI (mid, low-x), Nebula/Co-Star (low-y, mid-x), family astrologers (high-y, low-x, unscalable), ChatGPT (low-y, mid-x, free), Sitara (top-right, premium). Compare row: price/trust/AI/human/fulfilment/family. Takeaway: top-right quadrant is empty and defensible.

S27 • Our Right to Win — Only supported claims: AI product-building capability (this two-phase research + multi-agent workflow discipline is itself the demo); India + diaspora cultural fluency; language capability (Hinglish/Hindi product design); no legacy conflict (incumbents can't adopt ethics positioning without cannibalising fear-driven consult revenue — innovator's dilemma); founder commitment to trust-first brand. Honest gap: Jyotish credibility must be hired (named early-hire). Visual: strengths vs named-gaps table. Takeaway: our edge is product + trust posture; we hire the domain authority.

S28 • MVP & Validation Plan — MVP (12 wks, within ₹65L envelope): chart engine (licensed ephemeris) + daily brief (WhatsApp/web) + memory chat + Year-Ahead report + 5 contracted astrologers + billing. Not built: voice, community, apps, rituals. Validation gates: 20 interviews → landing test (>25% email capture) → 50-user concierge at \$15/mo (Gate: M2 renewal ≥60%, M3 ≥45%) → founding 500 at \$99 (Gate: CAC <₹4,000, activation >60%) → scale decision. Stop-loss: if concierge M2 <40% → pivot to membership+credits or stop; max loss capped at ₹1.25 Cr. Visual: stage-gate funnel with thresholds. Takeaway: every rupee buys a kill-or-scale answer.

S29 • GTM & Roadmap — First 100: r/ABCDesis + 2 diaspora micro-influencers + Honest-Kundli hook (organic-led). First 1,000: Meta lookalikes (AstroTalk's proven CAC ₹575-870 domestic benchmark; our NRI CAC guardrail ₹4,000 blended), festival campaigns, referral = gift a family chart. Channels later: SEO transit calendars, YouTube honest-astrology, wedding-platform partnerships (B2B2C). Roadmap: P1 Validate (M1-6) → P2 Productise (voice, family, M7-12) → P3 Scale (matchmaking reports, India tier, M13-24) → P4 Expand (rituals cross-sell = DharmaSetu module, M25-36) → P5 Platform (muhurat API, family-graph ecosystem). Visual: roadmap swimlane. Takeaway: validate narrow, expand along the family graph.

S30 • Board Recommendation — Approve: Scenario A — ₹1.25 Cr validation-and-build capital; milestone: concierge retention gate + 500 founding members by M6; next board review: month 6 with cohort data; if validation fails: stop-loss executed, ≤₹1.25 Cr total exposure, learnings + brand + panel reusable for pivot (KalaOS or hybrid-credits); if succeeds: Gate-2 ask ₹3.3 Cr for launch (Deck 2). Visual: decision box + timeline. Takeaway: a capped ₹65L to buy certainty on a ₹1,182-Cr-category premium niche.

DECK 2 — INVESTMENT, P&L, RETURNS & BUSINESS CASE (30 slides)

S1 • Cover — “Project Sitara — Investment Case & Three-Year Financial Plan.” Confidential. Note: all figures from bottom-up model (Excel attached); base case unless stated; ₹ lakh/crore; FX ₹90/\$.

S2 • Financial Executive Summary — Initial ₹1.25 Cr (Scenario A); 12-month total ₹4.55 Cr (A + B); monthly break-even ~M30 base (M23 aggressive / not-in-36-months conservative — stop-loss logic); revenue Y1 ₹1.42 Cr → Y2 ₹6.34 Cr → Y3 ₹12.16 Cr; Y3 EBITDA ≈ +₹0.20 Cr (turns

positive M30); peak cumulative cash need $\approx$ ₹7.85 Cr; 3-yr cumulative revenue $\approx$ ₹19.9 Cr; key risk: monthly churn (every +1pt $\approx$ -₹2.4 Cr Y3 revenue). Visual: stat cards + mini break-even chart.

S3 • Business Model (revenue streams) — Core model: monthly subs, annual subs (cash upfront, revenue recognised monthly), premium tier, add-on consult minutes, premium reports. Future (NOT in model): family plans priced separately, ritual fulfilment commission, B2B muhurat API, white-label. Visual: stream stack — core solid, future ghosted. Rule: nothing speculative in the P&L.

S4 • Pricing Architecture — Table (plan, target, monthly, annual, includes, human support, expected mix, plan GM): NRI Standard \$12.99/\$99 (45%); NRI Premium \$24.99/\$199 incl. 30 consult-min/qtr (12%); India Aarambh ₹499/₹3,999 (30%); India Pragati ₹999/₹7,999 incl. 15 min/qtr (13%). India prices GST-inclusive; model revenue net of GST (18%); NRI = zero-rated export. Blended net ARPU ₹815/mo + add-ons ₹65 $\rightarrow$ ₹880.

S5 • Core Assumptions (the honesty slide) — Full table with source/logic per line: CAC ₹3,600 (benchmark: AstroTalk ₹575-870 domestic; NRI premium assumed 4-5x); churn 6%/mo blended (benchmark: DTC 6.5% Recurly; annual-mix damped); visitor $\rightarrow$ trial 3%, trial $\rightarrow$ paid 25% (management estimates); referral 12% of new; refunds 2.5%; AI cost ₹70/user/mo; human-minutes ₹45; support ₹25; payment 2.8%; infra ₹15; FX ₹90. Colour-coded: fact/benchmark/assumption.

S6 • Customer Growth Model — 36-month table+chart: opening + new (mktg/CAC $\times$ 1.12 referral) - churn = closing. Milestones: M6 $\sim$ 900; M12 $\sim$ 3,200; M24 $\sim$ 8,600; M36 $\sim$ 17,400. Splits: monthly vs annual (60% annual), premium 25%, India 43%. Conservative M36 $\sim$ 8,000; aggressive $\sim$ 28,000.

S7 • Acquisition Funnel — Impressions 2.5M/mo $\rightarrow$ visitors 75K (3% CTR-adj) $\rightarrow$ leads 6,000 (8%) $\rightarrow$ trials 1,800 (3% of visitors) $\rightarrow$ paid 450 (25%) + referrals 54; cost/stage; funnel CAC math ties to ₹3,600.

S8 • CAC by Channel — Meta ₹4,200 · Google ₹4,800 · influencers ₹3,000 · community/organic ₹1,200 · referral ₹800 (chart-gift cost) · festival campaigns ₹3,400; blended ₹3,600. Benchmarks noted.

S9 • LTV & Ratios — LTV = ARPU ₹880 $\times$ GM 74% $\div$ churn 6% $\approx$ ₹10,850; LTV:CAC $\approx$ 3.0; CAC payback $\approx$ 5.5 months; with churn at 5% (Y3 target) LTV ₹13,000, ratio 3.6. Conservative case: 2.0 (marginal — the stop-loss logic). Chart: LTV vs CAC by scenario.

S10 • Unit Economics per Customer/Month — Waterfall: revenue ₹880 - AI -₹70 - human -₹45 - support -₹25 - payment -₹25 - infra -₹15 - refunds -₹22 $\rightarrow$ contribution ₹678 (74%) $\rightarrow$ CAC amortised (5.5-mo payback). Separate mini-tables: NRI Premium (best) vs India Aarambh (thinnest).

S11 • Investment Allocation (18-month, A+B ₹3.5 Cr) — Product & tech ₹1.32 Cr (38%); content & expert layer ₹42L (12%); marketing ₹1.05 Cr (30%); operations ₹28L (8%); legal & compliance ₹18L (5%); contingency ₹25L (7%). Sub-lines per your framework. Visual: allocation donut + phase split.

S12 • Scenario A — Lean Validation (₹65L) — 6 months; team 2 eng + founder + part-time Jyotish lead + designer-contract; 50 concierge users + 500 founding; expected revenue ₹32L (founding annuals, cash) ; success = M2 ≥60% concierge renewal, CAC <₹4,000; max acceptable loss ₹65L; decision gate month 6.

S13 • Scenario B — Focused Launch (+₹2.85 Cr) — M7-18; team to 9; marketing ₹15L/mo by M12; target 3,200 payers M12 / ~6,500 M18; revenue M7-18 ≈ ₹2.4 Cr; monthly burn peak ₹22L; monthly contribution break-even path visible by M18; milestones: churn ≤6%, blended CAC ≤₹3,600, NPS ≥55.

S14 • Scenario C — Accelerated Scale (+₹4.5-6.5 Cr at PMF) — M19-36; marketing to ₹35L/mo; India tier launch + matchmaking + rituals cross-sell; M36 ~17,400 payers, exit ARR ≈ ₹18.4 Cr; monthly EBITDA positive ~M27; risks: CAC inflation at scale, incumbent response; constraint stated: capital cannot buy retention — gates remain.

S15 • Year-1 Monthly P&L (table) — 12 columns: subs revenue (monthly+annual recognised), add-ons, COGS lines, GP, opex (salaries, marketing, tech, G&A, legal), EBITDA, cumulative burn. Y1: revenue ₹1.5 Cr, GP ₹1.1 Cr (74%), opex ₹3.4 Cr, EBITDA -₹2.3 Cr.

S16 • Y2-Y3 Annual P&L — Y2: revenue ₹6.2 Cr, GP ₹4.6 Cr, opex ₹6.8 Cr, EBITDA -₹2.2 Cr; Y3: revenue ₹13.7 Cr, GP ₹10.2 Cr, opex ₹9.0 Cr, EBITDA +₹1.2 Cr (9%); EBIT~EBITDA (asset-light, D&A ~₹8L); tax nil until loss carry-forwards absorbed (noted).

S17 • Cash Flow & Funding Plan — Opening + tranches (₹65L M1, ₹2.85 Cr M7, ₹4.5 Cr M19) – operating burn + annual-plan cash-ahead-of-revenue benefit (deferred revenue ≈ ₹1.1 Cr by M36); minimum cash floor ₹40L policy; peak cumulative need ₹6.4 Cr; closing cash M36 ≈ ₹1.9 Cr.

S18 • 36-Month Charts — Revenue ramp; customers; ARPU; churn glidepath 6→5%; GM%; marketing spend; monthly burn → EBITDA crossover (M27 flag).

S19 • Break-Even Analysis — Fixed cost M27 ≈ ₹52L/mo; contribution ₹656/user/mo (after refunds); BE customers ≈ 7,950; BE revenue ≈ ₹70L/mo; cash to BE ≈ ₹6.2 Cr. Conservative: BE not reached in 36 months → stop-loss/pivot rules; aggressive: M20, ~5,600 customers.

S20 • Scenario Comparison — 3-column: M36 customers 8,000/17,400/28,000; Y3 revenue ₹6.6/₹13.7/₹23.9 Cr; Y3 EBITDA -₹1.6/+₹1.2/+₹4.8 Cr; peak cash ₹5.1/₹6.4/₹8.9 Cr; BE month n.a./27/20. Chart: revenue fan.

S21 • Sensitivity (tornado) — Churn +25% (6→7.5%): Y3 revenue -₹2.9 Cr, BE +7 months; CAC +25%: peak cash +₹85L; ARPU -15%: Y3 EBITDA -₹1.5 Cr → negative; conversion -25%: M36 customers -24%; AI cost ×2: GM -8pts (74→66%); growth delayed 6 months: peak cash + ₹1.1 Cr; annual-mix 40% (vs 60%): deferred-cash benefit -₹45L; refunds ×2: GM -2.5pts. Top-3 drivers: churn, ARPU, CAC — all three are what validation measures.

S22 • Returns by Investment Level — Table for ₹50L / ₹1 Cr / ₹2.5 Cr / ₹5 Cr / ₹10 Cr (levels adjusted to the staged plan; user levels mapped): what it buys, team, runway, marketing, target payers, ARR, BE, 3-yr cum revenue/EBITDA, ROI range, constraint (explicit: capital doesn't buy retention; ₹10 Cr ≠ 2× the ₹5 Cr outcome — CAC inflation + gates). Includes “cash returned/retained” honesty row.

S23 · Idea-wise Financial Comparison — Top-3 (+DharmaSetu, Smriti) high-level models, same assumptions where fair: MVP invest, launch invest, price, GM%, CAC, churn, BE customers/month, Y1/Y3 revenue, Y3 EBITDA, 3-yr cash need, risk-adjusted score. Ranges, no false precision. Shows: KalaOS cheapest to BE (₹2.1 Cr, ~M22, but ₹9-11 Cr Y3 ceiling); Sahay biggest Y3 (₹16-22 Cr) at biggest cash (₹9-12 Cr) and ops risk; Sitara best balance.

S24 · What Could Kill This (downside case) — Conservative walk-through: churn 7.5%, CAC ₹4,500, ARPU ₹750 — M36 ~8,000 payers, cumulative EBITDA -₹6.3 Cr, no BE — and why we'd never get there: Gate-1 (M6) and Gate-2 (M12) stop capital at ₹65L / ₹2.1 Cr exposure. Stop-loss conditions enumerated.

S25 · Risk Register & Mitigation — Subscription-conversion risk (concierge gate); free-AI erosion (memory+humans+ethics moat); incumbent response (premium NRI niche + speed); regulatory/claims (ethics code, ASCI counsel); key-person Jyotish (codified corpus); FX; platform fees (web-first billing). Likelihood×impact heat map.

S26 · Governance & Reporting — Monthly MIS to board: payers, cohort retention curves, CAC by channel, NPS, cash; quarterly deep review; Gate calendar M6/M12/M18; independent cohort audit before Scenario C.

S27 · Board Decision Framework — Option 1 Do nothing (opportunity cost: category IPO window, niche gets claimed); Option 2 Validate only (₹65L, 6 months, capped risk — RECOMMENDED); Option 3 Build & launch now (₹3.5 Cr, skips cheapest learning, higher risk); Option 4 Partner/acquire (no suitable ethics-compatible target; AstroSage partnership optionality noted for Scenario C). Recommendation: Option 2.

S28 · The Ask — Approve ₹65L; use: product ₹28L, expert layer ₹8L, marketing ₹12L, ops/legal ₹9L, contingency ₹8L; 90-day deliverables: concierge cohort live, landing >25% capture, 20 interviews, chart engine alpha; 6-month target: 500 founding members, M2 retention ≥60%; 12-month target (post Gate-2): 3,200 payers, ₹1.5 Cr revenue run; stop-loss: concierge M2 <40% or founding <250 by M6; next funding decision: month 6 board.

S29 · Appendix pointers — Excel model map (17 sheets), formulas note, sources, rejected-idea file.

S30 · Thank You / Discussion — Contact + the one-line: “₹65 lakh buys the answer to a ₹1,182-crore question.”